

## Assignment 4

**Writeup-1:** Look at the assembly code above. The compiler has translated the code to set the start index at  $2^{16}$  and adds to it for each memory access. Why doesn't it set the start index to 0 and use small positive offsets?

这里的size定义为 $1L \ll 16$ 也就是 $2^{16}$ 。向量化是重写一个循环的过程，它不是对数组中的单个元素进行  $n$  次处理，而是同时对数组中的2个元素进行  $2^{16}/2$ 次处理，所以不是从index 0开始的。

**Write-up 2:** This code is still not aligned when using AVX2 registers. Fix the code to make sure it uses aligned moves for the best performance.

```
__builtin_assume_aligned(a, 16);
```

将这里的16改成32，因为AVX2 是 256 bits (也就是32 bytes) 的指令集。

**Write-up 3:** Provide a theory for why the compiler is generating dramatically different assembly.

`__builtin_assume_aligned(a, 16)`告诉编译器输入的C指针数据是16字节对齐的

MOVDQA - Move Aligned Double Quadword 将双四字从源操作数(第二个操作数)移动到目标操作数(第一个操作数)。

PMAXB - Maximum of Packed Unsigned Byte Integers 对目标操作数(第一个操作数)和源操作数(第二个操作数)中打包的无符号字节整数执行 SIMD 比较，并将每对字节整数的最大值返回到目标操作数。

由于 c 内存模型不允许 Clang 向内存位置引入额外的写操作。如果其他线程在调用上面的例程时修改 `a`，那么这些额外的写操作可能会影响修改最后的结果。这里使用了向量化的max指令，不需要使用任何显式的内部实现。最后汇编初的结果和之前也差距很大。

**Write-up 4:** Inspect the assembly and determine why the assembly does not include instructions with

vector registers. Do you think it would be faster if it did vectorize? Explain.

因为指针没有对齐，所以没法使用向量化寄存器。使用向量化的结果会更快。

向量化意味着一条指令可以并行地对多个操作数执行相同的操作。

假设正在处理的32位，单精度实数，这意味着一条指令可以一次执行8个操作，所以(至少在理论上)你可以完成  $n$  次赋值只用  $n/8$ 条乘法指令。至少在理论上，这应该允许操作以每次执行一条指令所允许的8倍的速度完成。

**Write-up 5:** Check the assembly and verify that it does in fact vectorize properly. Also what do you notice when you run the command

```
$ clang -O3 example4.c -o example4; ./example4
```

with and without the `-ffast-math` flag? Specifically, why do you see a difference in the

output.

```
wangwenqing@ubuntu:~/Desktop/A4-Code/recitation3$ clang -O3 example4.c -o example4; ./example4
The decimal floating point sum result is: 11.667578
The raw floating point sum result is: 0x1.755cccec10aa5p+3
wangwenqing@ubuntu:~/Desktop/A4-Code/recitation3$ clang -O3 -ffast-math example4.c -o example4; ./example4
The decimal floating point sum result is: 11.667578
The raw floating point sum result is: 0x1.755cccec10aa3p+3
wangwenqing@ubuntu:~/Desktop/A4-Code/recitation3$
```

查看汇编代码，已被正确向量化

当加了 -ffast-math 时，循环执行的顺序发生改变，也因此最后的sum result有微小的差距

**Write-up 6:** What speedup does the vectorized code achieve over the unvectorized code?

What additional speedup does using -mavx2 give? You may wish to run this experiment

several times and take median elapsed times; you can report answers to the nearest 100%

(e.g., 2×, 3×, etc). What can you infer about the bit width of the default vector registers on

the awsrn machines? What about the bit width of the AVX2 vector registers? *Hint:* aside

from speedup and the vectorization report, the most relevant information is that the data

type for each array is uint32\_t.

```
wangwenqing@ubuntu:~/Desktop/A4-Code/homework3$ make
clang -Wall -std=gnu99 -g -O3 -DNDEBUG -fno-vectorize -c loop.c
clang -o loop loop.o -lrt
wangwenqing@ubuntu:~/Desktop/A4-Code/homework3$ ls
fasttime.h  loop  loop.c  loop.o  Makefile
wangwenqing@ubuntu:~/Desktop/A4-Code/homework3$ time ./loop
Elapsed execution time: 0.054123 sec; N: 1024, I: 100000, __OP__: +, __TYPE__: u
int32_t

real    0m0.056s
user    0m0.043s
sys     0m0.004s
wangwenqing@ubuntu:~/Desktop/A4-Code/homework3$
```

make VECTORIZE=1 && time ./loop

```
wangwenqing@ubuntu:~/Desktop/A4-Code/homework3$ make VECTORIZE=1 && time ./loop
clang -Wall -std=gnu99 -g -O3 -DNDEBUG -Rpass=loop-vectorize -Rpass-missed=loop-
vectorize -ffast-math -c loop.c
loop.c:70:9: remark: vectorized loop (vectorization width: 4, interleaved count:
2) [-Rpass=loop-vectorize]
    for (j = 0; j < N; j++) {
    ^
clang -o loop loop.o -lrt
Elapsed execution time: 0.015824 sec; N: 1024, I: 100000, __OP__: +, __TYPE__: u
int32_t

real    0m0.021s
user    0m0.011s
sys     0m0.000s
wangwenqing@ubuntu:~/Desktop/A4-Code/homework3$
```

make VECTORIZE=1 AVX2=1 && time ./loop

```
wangwenqing@ubuntu:~/Desktop/A4-Code/homework3$ make VECTORIZE=1 AVX2=1 && time
./loop
clang -Wall -std=gnu99 -g -O3 -DNDEBUG -Rpass=loop-vectorize -Rpass-missed=loop-
vectorize -ffast-math -mavx2 -c loop.c
loop.c:70:9: remark: vectorized loop (vectorization width: 8, interleaved count:
4) [-Rpass=loop-vectorize]
    for (j = 0; j < N; j++) {
    ^
clang -o loop loop.o -lrt
Elapsed execution time: 0.009086 sec; N: 1024, I: 100000, __OP__: +, __TYPE__: u
int32_t

real    0m0.010s
user    0m0.006s
sys     0m0.000s
wangwenqing@ubuntu:~/Desktop/A4-Code/homework3$
```

vectorized 的速度比 unvectorized 快了4倍左右，因此其vector register 可能是128 bit

AVX2的速度 比 vectorized 快了2倍，因此其vector register 可能是256 bit